

CPE 342 - Assignment 2: Gradient Descent

Name	Student ID
Kiatisak Markmeeshap	67070501005

Fit $\hat{y} = c_0 + c_1 e^{c_2 x}$ to 100 observations by gradient descent.

```
In [1]: from pathlib import Path

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from scipy.optimize import curve_fit

pd.set_option("display.precision", 4)

df = pd.read_csv("data/gd_data.csv")
x = df["X"].to_numpy(dtype=float)
y = df["Y"].to_numpy(dtype=float)
m = x.size

print(f"m = {m} points")
print(f"x in [{x.min():.4f}, {x.max():.4f}]")
print(f"y in [{y.min():.4f}, {y.max():.4f}]")
df.head()
```

```
m = 100 points
x in [0.0108, 4.9783]
y in [0.9493, 1.9707]
```

```
Out[1]:
```

	X	Y
0	4.0723	0.9813
1	1.7243	1.1247
2	1.9020	1.0424
3	4.9141	0.9493
4	1.1659	1.1869

Task 1 - Loss function

Model:

$$\hat{y}_i = c_0 + c_1 e^{c_2 x_i}$$

Residual of the i -th point: $r_i = \hat{y}_i - y_i$.

Loss - the mean squared error over all m points:

$$L(c_0, c_1, c_2) = \frac{1}{m} \sum_{i=1}^m (\hat{y}_i - y_i)^2 = \frac{1}{m} \sum_{i=1}^m (c_0 + c_1 e^{c_2 x_i} - y_i)^2$$

```
In [2]: def predict(c, x):
        '''y_hat = c0 + c1 * exp(c2 * x)'''
        c0, c1, c2 = c
        return c0 + c1 * np.exp(c2 * x)

def loss(c, x, y):
    '''Mean squared error of the coefficients c on (x, y).'''
    return float(np.mean((predict(c, x) - y) ** 2))

C_INIT = np.zeros(3)
print(f"L(0, 0, 0) = {loss(C_INIT, x, y):.6f}")
```

L(0, 0, 0) = 1.305162

Task 2 - Gradient of the loss

Shared step. Each coefficient reaches L only through $\hat{y}_i$, so one chain rule covers all three:

$$\frac{\partial L}{\partial c_j} = \frac{1}{m} \sum_{i=1}^m \frac{\partial}{\partial c_j} (\hat{y}_i - y_i)^2 = \frac{1}{m} \sum_{i=1}^m 2 (\hat{y}_i - y_i) \frac{\partial \hat{y}_i}{\partial c_j} = \frac{2}{m} \sum_{i=1}^m r_i \frac{\partial \hat{y}_i}{\partial c_j}$$

Three branches. What is left is $\partial \hat{y}_i / \partial c_j$ for $\hat{y}_i = c_0 + c_1 e^{c_2 x_i}$.

$$\frac{\partial \hat{y}_i}{\partial c_0} = 1 \quad \frac{\partial \hat{y}_i}{\partial c_1} = e^{c_2 x_i} \quad \frac{\partial \hat{y}_i}{\partial c_2} = c_1 \frac{\partial}{\partial c_2} e^{c_2 x_i} = c_1 x_i e^{c_2 x_i}$$

The last one needs the chain rule a second time, on the exponent itself.

Result. Substituting each branch back into the shared step:

$$\frac{\partial L}{\partial c_0} = \frac{2}{m} \sum_{i=1}^m r_i$$

$$\frac{\partial L}{\partial c_1} = \frac{2}{m} \sum_{i=1}^m r_i e^{c_2 x_i}$$

$$\frac{\partial L}{\partial c_2} = \frac{2}{m} \sum_{i=1}^m r_i c_1 x_i e^{c_2 x_i}$$

where $r_i = c_0 + c_1 e^{c_2 x_i} - y_i$.

```
In [3]: def gradient(c, x, y):
        '''[dL/dc0, dL/dc1, dL/dc2] evaluated at c.'''
        c0, c1, c2 = c
        e = np.exp(c2 * x)
        r = c0 + c1 * e - y
        return (2 / x.size) * np.array([r.sum(), (r * e).sum(), (c1 * r * x *
        e).sum()])

print("grad L at (0.8, 1.2, -1.1) =", np.round(gradient([0.8, 1.2, -1.1], x,
y), 8))
```

grad L at (0.8, 1.2, -1.1) = [-0.25900113 -0.00712707 -0.02898728]

Task 3 - Update rule

With learning rate η , one gradient descent step is

$$\mathbf{c}^{(t+1)} = \mathbf{c}^{(t)} - \eta \nabla L(\mathbf{c}^{(t)})$$

which written out coefficient by coefficient is

$$c_0^{(t+1)} = c_0^{(t)} - \eta \frac{2}{m} \sum_{i=1}^m r_i$$

$$c_1^{(t+1)} = c_1^{(t)} - \eta \frac{2}{m} \sum_{i=1}^m r_i e^{c_2 x_i}$$

$$c_2^{(t+1)} = c_2^{(t)} - \eta \frac{2}{m} \sum_{i=1}^m r_i c_1 x_i e^{c_2 x_i}$$

All three residuals r_i are computed from $\mathbf{c}^{(t)}$, so the whole vector is updated at once, not one coefficient at a time.

```
In [4]: ETA = 0.3
        ITERATIONS = 5000
```

```

g0 = gradient(C_INIT, x, y)
print("c(0)      =", C_INIT)
print("grad L    =", np.round(g0, 6))
print("c(1)      =", np.round(C_INIT - ETA * g0, 6))

```

```

c(0)      = [0. 0. 0.]
grad L    = [-2.239027 -2.239027  0.      ]
c(1)      = [0.671708 0.671708 0.      ]

```

Note the third entry of the gradient: it carries a factor c_1 , so while $c_1 = 0$ the decay rate c_2 cannot move at all. The first step only lifts c_0 and c_1 off zero; c_2 starts moving from the second step onwards.

Task 4 - Implementation

Three routes to the same coefficients: an explicit loop that follows the formulas literally, a vectorised version, and `scipy.optimize.curve_fit` as an independent check.

Method 1 - explicit loop over the points

In [5]: `import math`

```

def gradient_descent_scalar(x, y, eta, iterations, c_init):
    '''Gradient descent written out point by point, straight from the
    formulas.'''
    c0, c1, c2 = c_init
    n = len(x)
    history = []
    for _ in range(iterations):
        g0 = g1 = g2 = 0.0
        squared = 0.0
        for xi, yi in zip(x, y):
            e = math.exp(c2 * xi)
            r = c0 + c1 * e - yi
            squared += r * r
            g0 += r
            g1 += r * e
            g2 += c1 * r * xi * e
        history.append(squared / n)
        c0 -= eta * 2 / n * g0
        c1 -= eta * 2 / n * g1
        c2 -= eta * 2 / n * g2
    return np.array([c0, c1, c2]), np.array(history)

c_scalar, history_scalar = gradient_descent_scalar(
    x.tolist(), y.tolist(), ETA, ITERATIONS, (0.0, 0.0, 0.0)

```

```
)  
print("c =", np.round(c_scalar, 6))
```

```
c = [ 1.003352  0.999271 -1.529872]
```

Method 2 - vectorised

```
In [6]: def gradient_descent(x, y, eta, iterations, c_init):  
        '''Same algorithm with numpy doing the sums. Returns coefficients and loss  
        history.'''  
        c = np.array(c_init, dtype=float)  
        history = np.empty(iterations)  
        for t in range(iterations):  
            history[t] = loss(c, x, y)  
            c = c - eta * gradient(c, x, y)  
        return c, history  
  
c_gd, history = gradient_descent(x, y, ETA, ITERATIONS, C_INIT)  
print("c =", np.round(c_gd, 6))
```

```
c = [ 1.003352  0.999271 -1.529872]
```

Method 3 - `scipy.optimize.curve_fit` (verification only)

```
In [7]: def exp_model(x, c0, c1, c2):  
        return c0 + c1 * np.exp(c2 * x)  
  
c_ref, _ = curve_fit(exp_model, x, y, p0=[1.0, 1.0, -1.0], maxfev=20000)  
print("c =", np.round(c_ref, 6))
```

```
c = [ 1.003353  0.999271 -1.529875]
```

Comparison

Methods 1 and 2 are the same arithmetic in a different shape, so they agree to the last bit. Method 3 lands a few units in the seventh decimal away: it uses a different algorithm (Levenberg-Marquardt) with its own stopping tolerance, so the gap measures where each one chose to stop, not a disagreement about the answer.

```
In [8]: comparison = pd.DataFrame({
    "Method 1 (loop)": [f"{v:.8f}" for v in c_scalar],
    "Method 2 (vectorised)": [f"{v:.8f}" for v in c_gd],
    "Method 3 (curve_fit)": [f"{v:.8f}" for v in c_ref],
    "|M1 - M2|": [f"{v:.2e}" for v in np.abs(c_scalar - c_gd)],
    "|M2 - M3|": [f"{v:.2e}" for v in np.abs(c_gd - c_ref)],
}, index=["c0", "c1", "c2"])

assert np.abs(c_scalar - c_gd).max() < 1e-12
assert np.abs(c_gd - c_ref).max() < 1e-5
comparison
```

```
Out[8]:
```

	Method 1 (loop)	Method 2 (vectorised)	Method 3 (curve_fit)	M1 - M2	M2 - M3
c0	1.00335243	1.00335243	1.00335257	0.00e+00	1.35e-07
c1	0.99927089	0.99927089	0.99927145	0.00e+00	5.56e-07
c2	-1.52987207	-1.52987207	-1.52987496	0.00e+00	2.89e-06

Task 5 - Results

Loss over the iterations

```
In [9]: checkpoints = [0, 1, 10, 100, 500, 1000, 2000, 3000, 4000, ITERATIONS - 1]
pd.DataFrame({
    "iteration": checkpoints,
    "loss (MSE)": [f"{history[k]:.8f}" for k in checkpoints],
}).set_index("iteration")
```

Out[9]:

loss (MSE)	
iteration	
0	1.30516214
1	0.10198385
10	0.01498087
100	0.00464958
500	0.00092328
1000	0.00076630
2000	0.00075549
3000	0.00075541
4000	0.00075541
4999	0.00075541

Resulting coefficients and fit quality

```
In [10]: c0, c1, c2 = c_gd
residual = y - predict(c_gd, x)

sse = float((residual ** 2).sum())
sst = float(((y - y.mean()) ** 2).sum())
mse = sse / m
rmse = float(np.sqrt(mse))
r_squared = 1 - sse / sst
grad_norm = float(np.linalg.norm(gradient(c_gd, x, y)))

print(f"c0 = {c0:.4f}")
print(f"c1 = {c1:.4f}")
print(f"c2 = {c2:.4f}")
print()
print(f"MSE = {mse:.6e}")
print(f"SSE = {sse:.4f}")
print(f"RMSE = {rmse:.4f}")
print(f"R^2 = {r_squared:.4f}")
print(f"||grad L|| after {ITERATIONS} iterations = {grad_norm:.2e}")

# A run that stalled somewhere else would fail here instead of printing junk.
assert grad_norm < 1e-7

print()
print(f"fitted model: y_hat = {c0:.4f} + {c1:.4f} * exp({c2:.4f} * x)")
```

```
c0 = 1.0034
c1 = 0.9993
c2 = -1.5299
```

```
MSE = 7.554135e-04
SSE = 0.0755
RMSE = 0.0275
R^2 = 0.9854
||grad L|| after 5000 iterations = 2.57e-08
```


fitted model: $y_{\text{hat}} = 1.0034 + 0.9993 * \exp(-1.5299 * x)$

```
In [11]: grid = np.linspace(x.min(), x.max(), 300)
fig, axes = plt.subplots(1, 2, figsize=(11, 4.2))

axes[0].scatter(x, y, s=18, alpha=0.7, label="data")
axes[0].plot(grid, predict(c_gd, grid), color="crimson", lw=2, label="fitted
curve")
axes[0].set(xlabel="x", ylabel="y", title="Fitted exponential model")
axes[0].legend()

axes[1].plot(history, lw=1.5)
axes[1].set(xlabel="iteration", ylabel="loss (MSE)", yscale="log",
            title=f"Loss over iterations (eta = {ETA})")

fig.tight_layout()
Path("figures").mkdir(exist_ok=True)
fig.savefig("figures/fit_and_loss.png", dpi=150, bbox_inches="tight")
plt.show()
```



Explanation

Write here.

► Numbers you can cite

- $c_0 = 1.0034$ - the level the curve settles at; at $x = 5$ the exponential term is only 0.0005
- $c_1 = 0.9993$ - the size of the decaying part; at $x = 0$ the model predicts $c_0 + c_1 = 2.0026$

- $c_2 = -1.5299$ - the decay rate; the gap to c_0 halves every $\ln 2 / 1.5299 = 0.4531$ in x
- fit quality: $\text{MSE} = 7.554 \times 10^{-4}$, $\text{RMSE} = 0.0275$, $R^2 = 0.9854$
- the sign of c_2 is what makes this decay rather than growth

Summary

```
In [12]: pd.DataFrame([
    ["1", "loss", "MSE", f"{mse:.6e}"],
    ["3", "update rule", "learning rate / iterations", f"{ETA} / {ITERATIONS}"],
    ["4", "coefficients", "c0, c1, c2", f"{c0:.4f}, {c1:.4f}, {c2:.4f}"],
    ["5", "fit", "SSE / RMSE / R^2", f"{sse:.4f} / {rmse:.4f} / {r_squared:.4f}"],
    ["5", "convergence", "||grad L|| at the end", f"{grad_norm:.2e}"],
], columns=["Task", "Quantity", "Measure", "Value"])
```

```
Out[12]:
```

	Task	Quantity	Measure	Value
0	1	loss	MSE	7.554135e-04
1	3	update rule	learning rate / iterations	0.3 / 5000
2	4	coefficients	c0, c1, c2	1.0034, 0.9993, -1.5299
3	5	fit	SSE / RMSE / R^2	0.0755 / 0.0275 / 0.9854
4	5	convergence	grad L at the end	2.57e-08

Closed form versus gradient descent

	Assignment 1 - OLS line	Assignment 2 - exponential model
Model	$\hat{y} = \beta_0 + \beta_1 x$	$\hat{y} = c_0 + c_1 e^{c_2 x}$
Closed form	yes - the normal equations	no - $\partial L / \partial c_2 = 0$ is transcendental
Learning rate	not needed	required, and the run fails if it is too large
Starting values	not needed	required; a poor start can reach a different answer
Global minimum	guaranteed - the loss is convex	not guaranteed - only a stationary point is reached
Cost	one matrix solve	thousands of passes over the data

The check that closes the gap: `curve_fit` reaches the same coefficients from a different algorithm, and the gradient norm printed above is of order $1e-8$ at the point gradient

descent stopped, so the run did settle at a genuine minimum.