

Activity - Estimating Lifetime

From `6_regression.ipynb` : build a Kaplan-Meier survival analysis on the Telco churn data and compare lifetime between payment methods.

survival term	Telco column
duration T	<code>tenure</code> , months as a customer
event E	<code>Churn == "Yes"</code>
censored	still a customer - we only know they lasted <i>at least</i> <code>tenure</code>
group	<code>PaymentMethod</code>

```
In [1]: from pathlib import Path

import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from lifelines import KaplanMeierFitter
from lifelines.statistics import multivariate_logrank_test, pairwise_logrank_test
from lifelines.utils import restricted_mean_survival_time

pd.set_option("display.precision", 3)

df = pd.read_csv("Telco-Churn.csv")
df["T"] = df["tenure"].astype(float)
df["E"] = (df["Churn"] == "Yes").astype(int)

print(df.shape)
df[["customerID", "PaymentMethod", "tenure", "Churn", "T", "E"]].head()
```

(7043, 23)

```
Out[1]:
```

	customerID	PaymentMethod	tenure	Churn	T	E
0	7590-VHVEG	Electronic check	1	No	1.0	0
1	5575-GNVDE	Mailed check	34	No	34.0	0
2	3668-QPYBK	Mailed check	2	Yes	2.0	1
3	7795-CFOCW	Bank transfer (automatic)	45	No	45.0	0
4	9237-HQITU	Electronic check	2	Yes	2.0	1

1. Censoring

The point of Kaplan-Meier: most rows are *censored*, not observed failures. Averaging `tenure` would treat "still here after 60 months" as a lifetime of 60, which understates every group.

```
In [2]: print(f"customers          {len(df)}")
print(f"churned (events)      {df.E.sum()}  ({df.E.mean():.1%})")
print(f"still active         {(1 - df.E).sum()}  (right-censored)")
print(f"tenure == 0           {(df['T'] == 0).sum()}  (brand new, no follow-up time)")
print()

# A zero-length observation carries no information for the estimator and lifelines
# warns about it, so those rows come out.
df = df[df["T"] > 0].copy()
print(f"kept {len(df)} rows")
```

```
customers          7043
churned (events)    1869  (26.5%)
still active        5174  (right-censored)
tenure == 0         11   (brand new, no follow-up time)
```


kept 7032 rows

2. Overall survival curve

```
In [3]: kmf = KaplanMeierFitter()
kmf.fit(df["T"], event_observed=df["E"], label="All customers")

ax = kmf.plot_survival_function(figsize=(7, 4.2))
ax.set(xlabel="tenure (months)", ylabel="P(still a customer)",
       title="Kaplan-Meier - all customers", ylim=(0, 1))
plt.tight_layout()
plt.show()

print("median survival time:", kmf.median_survival_time_)
print(kmf.survival_function_at_times([12, 24, 48, 72]).round(3))
```



```

median survival time: inf
12    0.843
24    0.789
48    0.709
72    0.593
Name: All customers, dtype: float64

```

3. By payment method

```


In [4]: METHODS = sorted(df["PaymentMethod"].unique())
fitters = {}

fig, ax = plt.subplots(figsize=(8, 5))
for method in METHODS:
    group = df[df["PaymentMethod"] == method]
    f = KaplanMeierFitter().fit(group["T"], event_observed=group["E"],
label=method)
    fitters[method] = f
    f.plot_survival_function(ax=ax, ci_show=True)

ax.set(xlabel="tenure (months)", ylabel="P(still a customer)",
       title="Customer survival by payment method", ylim=(0, 1))
ax.axhline(0.5, color="grey", lw=1, ls="--")
ax.legend(loc="lower left", fontsize=9)

plt.tight_layout()
Path("figures").mkdir(exist_ok=True)
fig.savefig("figures/km_by_payment_method.png", dpi=150, bbox_inches="tight")
plt.show()

```



4. Comparing lifetime

`median_survival_time_` is the statistic the lecture uses, but three of the four curves never reach 0.5 inside the 72 months observed, so it comes back as `inf` for them. Two statistics that do exist for every group:

- **survival at fixed times** - the share still subscribed at 12, 24, 48, 72 months
- **RMST** (restricted mean survival time) - the area under the curve up to 72 months, read as "expected months as a customer during the first 72"

In [5]: `HORIZON = 72`

```
rows = []
for method, f in fitters.items():
    s = f.survival_function_at_times([12, 24, 48, 72]).to_numpy()
    rows.append({
        "PaymentMethod": method,
        "n": int(f.event_observed.shape[0]),
        "churned": int(f.event_observed.sum()),
        "S(12)": s[0], "S(24)": s[1], "S(48)": s[2], "S(72)": s[3],
        "median": f.median_survival_time_,
        f"RMST({HORIZON})": restricted_mean_survival_time(f, t=HORIZON),
    })

summary =
pd.DataFrame(rows).set_index("PaymentMethod").sort_values(f"RMST({HORIZON})")
summary
```

Out[5]:

	n	churned	S(12)	S(24)	S(48)	S(72)	median	RMST(72)
PaymentMethod								
Electronic check	2365	1071	0.727	0.624	0.493	0.295	47.0	41.243
Mailed check	1604	308	0.828	0.797	0.740	0.726	inf	56.363
Bank transfer (automatic)	1542	258	0.931	0.902	0.847	0.745	inf	63.000
Credit card (automatic)	1521	232	0.945	0.913	0.855	0.758	inf	63.811

5. Are the differences real?

The log-rank test compares the curves without assuming any shape for them.

```
In [6]: overall = multivariate_logrank_test(df["T"], df["PaymentMethod"], df["E"])
print(f"all four groups: chi2 = {overall.test_statistic:.1f}, p = {overall.p_value:.3g}")
print()

pairwise = pairwise_logrank_test(df["T"], df["PaymentMethod"], df["E"])
pairwise.summary[["test_statistic", "p"]].round(4)
```

all four groups: chi2 = 865.2, p = 3.07e-187

Out[6]:

		test_statistic	p
Bank transfer (automatic)	Credit card (automatic)	0.868	0.351
	Electronic check	510.035	0.000
	Mailed check	51.072	0.000
Credit card (automatic)	Electronic check	539.740	0.000
	Mailed check	64.815	0.000
Electronic check	Mailed check	152.456	0.000

6. Cox proportional hazards

The curves above compare payment methods without holding anything else constant, and payment method is tangled with contract type: 78% of electronic-check customers are on month-to-month contracts against 36% of credit-card customers. A Cox model separates the two.

The coefficients are read as **hazard ratios**, `exp(coef)` : the multiplier on the instantaneous churn rate relative to the reference level, at every point in time. HR = 2 means twice the churn rate of the reference; HR = 1 means no difference.

Reference levels: credit card (automatic), two-year contract.

```
In [7]: from lifelines import CoxPHFitter
from lifelines.statistics import proportional_hazard_test

REFERENCE = ["PaymentMethod_Credit card (automatic)", "Contract_Two year"]

design = pd.get_dummies(df[["PaymentMethod", "Contract"]],
dtype=float).drop(columns=REFERENCE)
design["tenure"] = df["T"].to_numpy()
design["churned"] = df["E"].to_numpy()

payment_cols = [c for c in design.columns if c.startswith("PaymentMethod")]
design.head()
```

```
Out[7]:
```

	PaymentMethod_Bank transfer (automatic)	PaymentMethod_Electronic check	PaymentMethod_Mailed check	Contract_1 to-
0	0.0	1.0	0.0	
1	0.0	0.0	1.0	
2	0.0	0.0	1.0	
3	1.0	0.0	0.0	
4	0.0	1.0	0.0	

Model A - payment method alone

```
In [8]: cox_a = CoxPHFitter().fit(design[payment_cols + ["tenure", "churned"]],
"tenure", "churned")

cox_a.summary[["exp(coef)", "exp(coef) lower 95%", "exp(coef) upper 95%", "p"]]
```

Out[8]:

	exp(coef)	exp(coef) lower 95%	exp(coef) upper 95%	p
covariate				
PaymentMethod_Bank transfer (automatic)	1.089	0.912	1.301	3.446e-01
PaymentMethod_Electronic check	4.621	4.005	5.332	1.130e-97
PaymentMethod_Mailed check	2.149	1.810	2.551	2.540e-18

Model B - payment method adjusted for contract type

```
In [9]: cox_b = CoxPHFitter().fit(design, "tenure", "churned")

print(f"concordance model A {cox_a.concordance_index_:.3f} model B  
{cox_b.concordance_index_:.3f}")
cox_b.summary[["exp(coef)", "exp(coef) lower 95%", "exp(coef) upper 95%", "p"]]
```

concordance model A 0.676 model B 0.809

Out[9]:


	exp(coef)	exp(coef) lower 95%	exp(coef) upper 95%	p
covariate				
PaymentMethod_Bank transfer (automatic)	1.028	0.861	1.228	7.593e-01
PaymentMethod_Electronic check	2.104	1.822	2.430	3.981e-24
PaymentMethod_Mailed check	1.848	1.555	2.197	3.447e-12
Contract_Month-to-month	54.102	39.712	73.706	3.566e-141
Contract_One year	6.888	4.944	9.598	4.068e-30

```
In [10]: comparison = pd.DataFrame({
    "HR alone (A)": cox_a.hazard_ratios_,
    "HR adjusted (B)": cox_b.hazard_ratios_[payment_cols],
})
comparison["shrinkage"] = 1 - comparison["HR adjusted (B)"] / comparison["HR  
alone (A)"]
comparison
```

Out[10]:

	HR alone (A)	HR adjusted (B)	shrinkage
covariate			
PaymentMethod_Bank transfer (automatic)	1.089	1.028	0.056
PaymentMethod_Electronic check	4.621	2.104	0.545
PaymentMethod_Mailed check	2.149	1.848	0.140

```
In [11]: ax = cox_b.plot()
ax.set_title("Model B - log hazard ratios with 95% CI")
plt.tight_layout()
plt.savefig("figures/cox_hazard_ratios.png", dpi=150, bbox_inches="tight")
plt.show()
```



Does the model's own assumption hold?

Cox assumes a covariate's hazard ratio is *constant over time*. That is testable, and here it fails for some covariates - so each HR above should be read as an average effect across the 72 months, not a fixed multiplier. The mailed-check curve in section 3 shows why: it drops steeply in the first months and then flattens.

```
In [12]: proportional_hazard_test(cox_b, design,
time_transform="rank").summary[["test_statistic", "p"]]
```

Out[12]:

	test_statistic	p
Contract_Month-to-month	20.686	5.411e-06
Contract_One year	0.008	9.286e-01
PaymentMethod_Bank transfer (automatic)	0.303	5.821e-01
PaymentMethod_Electronic check	9.082	2.582e-03
PaymentMethod_Mailed check	30.595	3.180e-08

7. Minimum follow-up time

Two different questions hide under "how long do we need to watch customers".

How soon can a conclusion be drawn? Censor every customer administratively at month `t` - as if the study had stopped there - and re-run the log-rank test. The smallest `t` whose test comes back significant is the minimum follow-up that particular comparison needs.

How long do the estimates stay trustworthy? A survival curve is only as good as the number of customers still at risk behind it. Once the risk set thins out, the curve is carried by a handful of people and the confidence band widens accordingly.

```
In [13]: HORIZONS = range(1, 25)
T_all, E_all, G_all = df["T"].to_numpy(), df["E"].to_numpy(),
df["PaymentMethod"].to_numpy()

first_significant = {}
for t in HORIZONS:
    censored_T = np.minimum(T_all, t)
    censored_E = ((E_all == 1) & (T_all <= t)).astype(int)
    for pair, p in pairwise_logrank_test(censored_T, G_all,
censored_E).summary["p"].items():
        if pair not in first_significant and p < 0.05:
            first_significant[pair] = (t, p)

full = pairwise_logrank_test(T_all, G_all, E_all).summary["p"]
rows = []
for pair in full.index:
    t, p = first_significant.get(pair, (np.nan, np.nan))
    rows.append({"pair": " vs ".join(pair),
                "min follow-up (months)": t,
                "p at that point": p,
                "p on full 72 months": full[pair]})

pd.DataFrame(rows).sort_values("min follow-up (months)").set_index("pair")
```


Out[13]:

	min follow-up (months)	p at that point	p on full 72 months
pair			
Bank transfer (automatic) vs Electronic check	1.0	5.673e-25	6.231e-113
Bank transfer (automatic) vs Mailed check	1.0	1.930e-19	8.905e-13
Credit card (automatic) vs Electronic check	1.0	1.765e-28	2.148e-119
Credit card (automatic) vs Mailed check	1.0	4.866e-23	8.226e-16
Electronic check vs Mailed check	2.0	5.579e-03	5.038e-35
Bank transfer (automatic) vs Credit card (automatic)	NaN	NaN	3.514e-01

```
In [14]: at_risk = pd.DataFrame({
    method: [(group["T"] >= t).mean() for t in (12, 24, 36, 48, 60, 72)]
    for method, group in df.groupby("PaymentMethod")
}, index=[f"t={t}" for t in (12, 24, 36, 48, 60, 72)].T
at_risk.columns.name = "fraction still at risk"

fig, ax = plt.subplots(figsize=(7.5, 4.2))
for method, group in df.groupby("PaymentMethod"):
    months = np.arange(0, 73)
    ax.plot(months, [(group["T"] >= t).mean() for t in months], lw=1.5,
label=method)
ax.axhline(0.10, color="grey", ls="--", lw=1)
ax.set(xlabel="tenure (months)", ylabel="fraction still at risk",
    title="Risk set thinning - below the dashed line the tail is guesswork",
ylim=(0, 1))
ax.legend(fontsize=9)
plt.tight_layout()
plt.savefig("figures/at_risk_thinning.png", dpi=150, bbox_inches="tight")
plt.show()

at_risk
```



Out[14]:

	fraction still at risk	t=12	t=24	t=36	t=48	t=60	t=72
Bank transfer (automatic)		0.871	0.748	0.627	0.507	0.355	0.104
Credit card (automatic)		0.863	0.748	0.621	0.500	0.347	0.095
Electronic check		0.605	0.440	0.311	0.210	0.113	0.017
Mailed check		0.553	0.372	0.252	0.165	0.087	0.011

8. Conclusion

```
In [15]: best, worst = summary.index[-1], summary.index[0]
gap = summary.loc[best, f"RMST({HORIZON})"] - summary.loc[worst,
f"RMST({HORIZON})"]

print(f"longest-lived group {best} RMST {summary.loc[best,
f'RMST({HORIZON})']:.1f} months")
print(f"shortest-lived group {worst} RMST {summary.loc[worst,
f'RMST({HORIZON})']:.1f} months")
print(f"gap {gap:.1f} months over a 72-month horizon")
print()
print(f"electronic-check hazard ratio, unadjusted
{cox_a.hazard_ratios_['PaymentMethod_Electronic check']:.2f}")
print(f"electronic-check hazard ratio, given contract
{cox_b.hazard_ratios_['PaymentMethod_Electronic check']:.2f}")
```

longest-lived group Credit card (automatic) RMST 63.8 months
shortest-lived group Electronic check RMST 41.2 months
gap 22.6 months over a 72-month horizon

electronic-check hazard ratio, unadjusted	4.62
electronic-check hazard ratio, given contract	2.10

Lifetime by payment method. Automatic payment lasts longest - about 63 of the first 72 months as a customer - and the two automatic methods cannot be told apart at all: the log-rank test gives $p = 0.35$ on the full data, and it never falls below 0.05 at any follow-up horizon. Mailed check sits in the middle at 56 months, electronic check is far behind at 41. Median survival time, the statistic the lecture uses, only exists for electronic check (47 months); the other three curves never fall to 0.5, so RMST does the comparing here.

How much of that is the payment method itself. On its own, electronic check carries a hazard ratio of 4.6 against credit card. Once contract type is held constant it falls to about 2.1 - so roughly half the apparent effect was contract type wearing a payment method's clothes. What is left is still a doubled churn rate, and mailed check behaves much the same way once adjusted, which points at the manual methods themselves rather than at who chooses them.

Minimum follow-up. Every comparison the data can support is already significant after **2 months** of follow-up: four of the five pairs separate within the first month, and electronic check versus mailed check needs a second. Watching longer sharpens the estimates but changes no conclusion - except for automatic bank transfer versus credit card, which stays indistinguishable however long anyone waits. At the other end, fewer than 10% of the manual-payment customers are still at risk past 60 months, so the tail of those curves should not be quoted.